

DSBDA Lab – Viva Questions & Answers

Marathwada Mitra Mandal's Institute of Technology (MMIT), Pune

T.E. Computer Engineering | 310256: DSBDA Lab | 2025–26

SPPU-Style – Simple & Clear Q&A

Group A (Python – Data Science): Practicals 1, 2, 3, 4, 5, 6, 7, 9, 10 – 10 questions each

Group B (Java/Scala – Big Data): Practicals B1 (MapReduce Word Count), B2 (Log File Processing), B3 (Apache Spark/Scala) – 10 questions each

GROUP A – Data Science (Python)

Practical 1 – Data Wrangling I

Aim: Perform data wrangling on an open-source dataset: load data, handle missing values, normalize and encode categorical data using Python.

Q.No	Question	Answer
1	What is Data Wrangling?	Data wrangling is the process of cleaning, transforming and organizing raw data into a usable format for analysis. It includes handling missing values, removing duplicates, formatting and encoding.
2	What is a DataFrame in Pandas?	A DataFrame is a 2-dimensional labelled data structure in Pandas, like a table with rows and columns. It is created using <code>pd.read_csv()</code> or <code>pd.DataFrame()</code> .
3	How do you check for missing values in Pandas?	Using <code>df.isnull().sum()</code> – it returns the count of missing (NaN) values in each column.
4	What is the difference between <code>dropna()</code> and <code>fillna()</code> ?	<code>dropna()</code> removes rows or columns with missing values. <code>fillna()</code> replaces missing values with a given value like mean, median or 0 instead of removing them.
5	What is normalization? Give one formula.	Normalization scales data to a standard range. Min-Max normalization formula: $X_{\text{new}} = (X - X_{\text{min}}) / (X_{\text{max}} - X_{\text{min}})$. This scales values between 0 and 1.
6	What is the difference between normalization and standardization?	Normalization scales data to [0,1] range. Standardization converts data to have mean=0 and std=1 using Z-score: $Z = (X - \text{mean}) / \text{std}$.
7	What is Label Encoding?	Label encoding converts categorical text values into numbers. For example, Male=0, Female=1. It is done using <code>LabelEncoder</code> from <code>sklearn.preprocessing</code> .
8	What is One-Hot Encoding?	One-hot encoding creates a new binary column for each category. For example, a 'color' column with Red, Blue, Green becomes three separate 0/1 columns.
9	What does <code>df.describe()</code> do?	<code>df.describe()</code> gives a statistical summary of numeric columns including count, mean, std, min, 25%, 50%, 75% and max values.
10	What is the use of <code>df.info()</code> ?	<code>df.info()</code> shows the number of rows, column names, data types (dtype) and non-null counts. It helps understand the structure of the dataset.

Practical 2 – Data Wrangling II

Aim: *Handle missing values, detect outliers using boxplots/IQR, and apply data transformations on the Academic Performance dataset.*

Q.No	Question	Answer
1	What is an outlier?	An outlier is a data point that is significantly different from other observations in the dataset. It can be unusually high or low compared to the rest of the data.
2	How do you detect outliers using IQR?	$IQR = Q3 - Q1$. Outliers are values below $Q1 - 1.5 * IQR$ or above $Q3 + 1.5 * IQR$. $Q1$ is 25th percentile and $Q3$ is 75th percentile.
3	What is imputation?	Imputation is the process of replacing missing values with estimated values such as mean, median, or mode of the column instead of deleting the row.
4	When do you use mean vs median imputation?	Use mean when data is symmetric (no outliers). Use median when data is skewed or has outliers, as median is not affected by extreme values.
5	What is a boxplot and what does it show?	A boxplot shows the distribution of data using five values: minimum, $Q1$, median ($Q2$), $Q3$ and maximum. Points outside the whiskers are outliers.
6	What is data transformation?	Data transformation changes the scale or distribution of data. Common types include log transformation, square root and Box-Cox transformation to make data more normally distributed.
7	What is log transformation? When is it used?	Log transformation applies $\log(x)$ to each value. It is used when data is right-skewed (positively skewed) to reduce the effect of large values.
8	What is the difference between <code>drop()</code> and <code>dropna()</code> ?	<code>drop()</code> removes specific rows or columns by label. <code>dropna()</code> removes rows/columns that contain missing (NaN) values.
9	What is <code>df.duplicated()</code> used for?	<code>df.duplicated()</code> returns True for rows that are duplicates. <code>df.drop_duplicates()</code> removes those duplicate rows from the DataFrame.
10	What Python library is commonly used for data wrangling?	Pandas is the most commonly used library for data wrangling in Python. NumPy is used for numerical operations. Matplotlib and Seaborn are used for visualization.

Practical 3 – Descriptive Statistics

Aim: Compute measures of central tendency (mean, median, mode) and variability (std, variance, percentiles) on the Iris dataset grouped by species.

Q.No	Question	Answer
1	What is descriptive statistics?	Descriptive statistics summarises and describes the main features of a dataset. It includes measures of central tendency (mean, median, mode) and measures of spread (variance, std).
2	What is mean? How is it calculated?	Mean is the average of all values. It is calculated as: Mean = Sum of all values / Total number of values.
3	What is the difference between mean and median?	Mean is the average of all values. Median is the middle value when data is sorted. Median is preferred when data has outliers as mean gets affected by extreme values.
4	What is standard deviation?	Standard deviation measures how spread out data values are from the mean. A small std means data is close to the mean; a large std means data is widely spread.
5	What is variance?	Variance is the average of squared differences from the mean. It is the square of standard deviation. Formula: Variance = std^2 .
6	What is a percentile?	A percentile is the value below which a given percentage of data falls. For example, the 75th percentile means 75% of values are below that value.
7	What does groupby() do in Pandas?	groupby() groups data based on a column value and then applies an aggregation function like mean(), sum(), count() or describe() on each group.
8	What is the Iris dataset?	The Iris dataset contains 150 samples of flowers with 4 features: sepal length, sepal width, petal length, petal width. It has 3 species: setosa, versicolor, virginica.
9	What is the difference between df.mean() and df.groupby().mean()?	df.mean() gives the mean of all values in each column. df.groupby('col').mean() gives the mean for each group separately, like mean per species.
10	What is skewness?	Skewness measures the asymmetry of a distribution. Positive skew means tail is on the right. Negative skew means tail is on the left. Zero skew means symmetric (normal) distribution.

Practical 4 – Linear Regression (Boston Housing)

Aim: Build a Linear Regression model to predict house prices using the Boston Housing dataset with 506 samples and 14 features.

Q.N o	Question	Answer
1	What is Linear Regression?	Linear Regression is a supervised machine learning algorithm that predicts a continuous output variable based on one or more input variables using a straight-line relationship: $Y = mX + c$.
2	What is the difference between Simple and Multiple Linear Regression?	Simple Linear Regression has one input feature ($Y = mX + c$). Multiple Linear Regression has two or more input features ($Y = b_0 + b_1X_1 + b_2X_2 + \dots$).
3	What is the Boston Housing dataset?	It is a dataset with 506 samples and 13 input features (like CRIM, RM, LSTAT) used to predict MEDV – the median house price. It is a standard regression benchmark dataset.
4	What is train-test split? Why is it done?	Train-test split divides the dataset into training data (to train the model) and test data (to evaluate performance). Commonly split as 80% train and 20% test.
5	What is Mean Squared Error (MSE)?	MSE measures the average squared difference between actual and predicted values. Lower MSE means better model. Formula: $MSE = (1/n) * \sum (y_{actual} - y_{predicted})^2$.
6	What is R^2 score (coefficient of determination)?	R^2 score measures how well the model fits the data. Value ranges from 0 to 1. $R^2=1$ means perfect prediction. $R^2=0$ means the model is no better than just predicting the mean.
7	What does the slope (coefficient) in Linear Regression represent?	The slope (m or coefficient) represents the change in the output variable for a one-unit increase in the input variable, keeping all other features constant.
8	Which library is used for Linear Regression in Python?	from sklearn.linear_model import LinearRegression. Then create model = LinearRegression(), fit with model.fit(X_train, y_train), predict with model.predict(X_test).
9	What is overfitting?	Overfitting occurs when a model performs very well on training data but poorly on test data. It means the model has memorised the training data instead of learning general patterns.

10 What is the difference between regression and classification?

Regression predicts a continuous numeric output (e.g., house price). Classification predicts a discrete class label (e.g., spam or not spam).

Practical 5 – Logistic Regression (Social Network Ads)

Aim: Implement Logistic Regression on Social_Network_Ads.csv for binary classification and compute Confusion Matrix metrics.

Q.No	Question	Answer
1	What is Logistic Regression?	Logistic Regression is a supervised classification algorithm used for binary output (0 or 1). It uses the sigmoid function to predict the probability of a class.
2	What is the sigmoid function?	Sigmoid function: $\sigma(z) = 1 / (1 + e^{(-z)})$. It maps any real value to a range between 0 and 1, which can be interpreted as probability.
3	What is a Confusion Matrix?	A confusion matrix is a table showing TP (True Positive), FP (False Positive), TN (True Negative) and FN (False Negative) to evaluate classification performance.
4	What is Accuracy?	Accuracy = $(TP + TN) / (TP + TN + FP + FN)$. It tells what percentage of total predictions were correct.
5	What is Precision?	Precision = $TP / (TP + FP)$. It measures how many of the predicted positive results were actually positive. High precision means fewer false alarms.
6	What is Recall (Sensitivity)?	Recall = $TP / (TP + FN)$. It measures how many of the actual positives were correctly identified. High recall means fewer actual positives were missed.
7	What is the difference between Logistic and Linear Regression?	Linear Regression predicts continuous values. Logistic Regression predicts class labels (0 or 1) for classification using the sigmoid function.
8	What is the Social Network Ads dataset?	It contains user data like Age and EstimatedSalary. The task is to predict whether a user will purchase a product (1) or not (0) – a binary classification problem.
9	What is the Error Rate?	Error Rate = $(FP + FN) / (TP + TN + FP + FN) = 1 - \text{Accuracy}$. It tells what fraction of all predictions were wrong.
10	What is the F1-score?	F1-score = $2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$. It is the harmonic mean of precision and recall, useful when there is class imbalance.

Practical 6 – Naïve Bayes Classifier (Iris Dataset)

Aim: Implement the Naïve Bayes classifier on the Iris dataset and compute Confusion Matrix metrics.

Q.No	Question	Answer
1	What is Naïve Bayes classifier?	Naïve Bayes is a probabilistic classification algorithm based on Bayes' Theorem. It assumes all features are independent of each other (naïve assumption).
2	State Bayes' Theorem.	$P(C X) = P(X C) \times P(C) / P(X)$. Where $P(C X)$ is posterior probability, $P(X C)$ is likelihood, $P(C)$ is prior probability and $P(X)$ is evidence.
3	What does 'naïve' mean in Naïve Bayes?	Naïve means that the algorithm assumes all input features are independent of each other. In reality this may not be true, but the algorithm still works well in practice.
4	What is GaussianNB and why is it used for Iris?	GaussianNB assumes features follow a Gaussian (normal) distribution. It is used for Iris because the features (sepal/petal length, width) are continuous numeric values.
5	What is prior probability?	Prior probability $P(C)$ is the probability of a class before observing any data. For Iris with equal samples: $P(\text{setosa}) = P(\text{versicolor}) = P(\text{virginica}) = 1/3 = 0.33$.
6	What is posterior probability?	Posterior probability $P(C X)$ is the updated probability of a class after observing the feature values. Naïve Bayes predicts the class with highest posterior probability.
7	What is the difference between GaussianNB, MultinomialNB and BernoulliNB?	GaussianNB is for continuous features. MultinomialNB is for word count/frequency data (text). BernoulliNB is for binary features (present/absent).
8	Why is Naïve Bayes fast?	Naïve Bayes is fast because it only needs to calculate class probabilities and feature probabilities from training data. There is no complex optimization or iteration involved.
9	What is Laplace Smoothing?	Laplace smoothing adds a small value (usually 1) to all feature counts to avoid zero probability for features not seen in training data. It prevents the zero-probability problem.

10 What are advantages of Naïve Bayes?

Advantages: (1) Simple and fast. (2) Works well with small datasets. (3) Handles both binary and multi-class classification. (4) Not sensitive to irrelevant features.

Practical 7 – Text Analytics & TF-IDF

Aim: Perform document preprocessing: Tokenization, POS Tagging, Stopword Removal, Stemming, Lemmatization and compute TF-IDF.

Q.No	Question	Answer
1	What is Text Analytics?	Text Analytics (also called Text Mining) is the process of extracting useful information from unstructured text data. It includes tokenization, POS tagging, stemming, lemmatization, etc.
2	What is Tokenization?	Tokenization is the process of splitting a sentence or paragraph into smaller units called tokens (words or sentences). NLTK provides <code>word_tokenize()</code> and <code>sent_tokenize()</code> .
3	What are Stopwords?	Stopwords are common words like 'is', 'the', 'a', 'an', 'in' that carry little meaning. They are removed before text analysis to reduce noise. NLTK provides a built-in stopwords list.
4	What is Stemming?	Stemming reduces words to their root form by removing suffixes using rules. Example: 'running' → 'run', 'studies' → 'studi'. Porter Stemmer is commonly used in NLTK.
5	What is Lemmatization?	Lemmatization converts a word to its base dictionary form (lemma) using vocabulary. Example: 'studies' → 'study', 'running' → 'run'. It is more accurate than stemming.
6	What is the difference between stemming and lemmatization?	Stemming is faster but may produce non-real words (e.g., 'studi'). Lemmatization is slower but always produces a real dictionary word (e.g., 'study').
7	What is POS Tagging?	POS (Part-of-Speech) tagging assigns grammatical tags to words like Noun (NN), Verb (VB), Adjective (JJ). NLTK's <code>pos_tag()</code> function is used for this.
8	What is TF-IDF?	TF-IDF stands for Term Frequency–Inverse Document Frequency. It gives high weight to words that appear frequently in a document but rarely across all documents, showing their importance.
9	What is Term Frequency (TF)?	$TF = (\text{Number of times word appears in document}) / (\text{Total words in document})$. It measures how often a word appears in a specific document.

10 What is Inverse Document Frequency (IDF)?

IDF = $\log(\text{Total number of documents} / \text{Number of documents containing the word})$. Words appearing in many documents get a lower IDF score (less important).

Practical 9 – Data Visualization II (Titanic – Box Plot)

Aim: Use Seaborn's inbuilt Titanic dataset to plot a box plot showing age distribution by gender and survival status.

Q.No	Question	Answer
1	What is Data Visualization?	Data visualization is the graphical representation of data and information using charts, graphs and plots. It helps to understand patterns, trends and outliers in data easily.
2	What is a Box Plot?	A box plot displays the distribution of data using five values: minimum, Q1, median, Q3 and maximum. Points outside the whiskers are outliers. It is used to compare distributions.
3	What does the box in a box plot represent?	The box spans from Q1 (25th percentile) to Q3 (75th percentile). The line inside the box represents the median (Q2 or 50th percentile). The box size = IQR.
4	What are whiskers in a box plot?	Whiskers extend from Q1 to the minimum and from Q3 to the maximum, within $1.5 \times \text{IQR}$ range. Points beyond the whiskers are plotted as individual dots and considered outliers.
5	What is the Titanic dataset?	The Titanic dataset contains information about passengers of the Titanic ship including age, sex, pclass, fare, sibsp, parch and survived (0=no, 1=yes). It is available in Seaborn.
6	How do you load the Titanic dataset in Seaborn?	<code>import seaborn as sns; titanic = sns.load_dataset('titanic')</code> . This directly loads the Titanic DataFrame without needing any CSV file.
7	What is the hue parameter in Seaborn?	hue parameter adds a third dimension to the plot by colour-coding data points based on another column. For example, <code>hue='survived'</code> shows survival status in different colours.
8	What is the difference between a box plot and a histogram?	A histogram shows frequency distribution of a single variable. A box plot shows quartiles, median and outliers and is better for comparing distributions across different groups.
9	What is <code>seaborn.boxplot()</code> syntax?	<code>sns.boxplot(x='sex', y='age', hue='survived', data=titanic)</code> . This plots age distribution for each gender, further split by survival status.

10	What observation can be made from Titanic age vs survival box plot?	Young female passengers had a higher survival rate. Children generally survived more. Older males had lower survival rates. This shows the 'women and children first' evacuation pattern.
----	---	---

Practical 10 – Data Visualization III (Iris – Histogram & Boxplot)

Aim: Load the Iris dataset, scan feature types, create histograms and box plots for each feature, compare distributions and identify outliers.

Q.No	Question	Answer
1	What is a Histogram?	A histogram is a bar chart that shows the frequency distribution of a continuous variable. The x-axis shows value ranges (bins) and y-axis shows the count of values in each bin.
2	How do you plot a histogram in Python?	Using matplotlib: <code>plt.hist(df['column'], bins=10)</code> . Using Seaborn: <code>sns.histplot(df['column'], bins=10)</code> . bins controls the number of intervals.
3	What is the difference between a histogram and a bar chart?	A histogram is used for continuous data with no gaps between bars. A bar chart is used for categorical data with gaps between bars.
4	How many features does the Iris dataset have? Name them.	The Iris dataset has 4 numeric features: sepal_length, sepal_width, petal_length, petal_width. It also has 1 categorical feature: species (setosa, versicolor, virginica).
5	What is df.dtypes used for?	df.dtypes returns the data type of each column – like int64, float64, or object. This helps identify which features are numeric (quantitative) and which are categorical (nominal).
6	What is the use of plt.subplot() in matplotlib?	plt.subplot(rows, cols, index) creates multiple plots in a single figure. For example, plt.subplot(2, 2, 1) creates a 2x2 grid of plots and places the plot at position 1.
7	How do you plot a boxplot for all features in Iris?	df.boxplot() or sns.boxplot(data=df) plots box plots for all numeric columns at once. You can also use a loop: <code>for col in df.columns: sns.boxplot(y=df[col])</code> .
8	What is the shape of the Iris dataset?	Iris has 150 rows and 5 columns (4 numeric features + 1 species column). Each of the 3 species has exactly 50 samples.
9	What observation can be made from Iris petal_length histogram?	The petal_length histogram shows a bimodal distribution – one peak around 1-2 cm (Iris setosa) and another around 4-6 cm (Iris versicolor and virginica).

10	Which Seaborn function shows both histogram and KDE curve?	<code>sns.histplot(df['column'], kde=True)</code> plots a histogram with a KDE (Kernel Density Estimate) curve overlaid, showing the smooth probability distribution of the data.
----	--	---

GROUP B – Big Data Analytics (Java / Scala)

Group B – Practical 1: Hadoop MapReduce Word Count (Java)

Aim: Write a Java program for Word Count using the Hadoop MapReduce framework on local standalone setup.

Q.No	Question	Answer
1	What is Hadoop?	Hadoop is an open-source framework for storing and processing large datasets in a distributed environment across clusters of computers. It has two main components: HDFS (storage) and MapReduce (processing).
2	What is MapReduce?	MapReduce is a programming model for processing large datasets in parallel. It has two phases: Map phase (splits data and processes it) and Reduce phase (aggregates results from mappers).
3	What is the Map phase in Word Count?	In the Map phase, each line of text is split into words. For each word, the mapper emits a key-value pair: (word, 1). Example: 'Hello World' → (Hello,1), (World,1).
4	What is the Reduce phase in Word Count?	In the Reduce phase, the reducer receives all values for a key and sums them up. Example: (Hello, [1,1,1]) → (Hello, 3). It gives the final word count.
5	What is HDFS?	HDFS (Hadoop Distributed File System) is the storage component of Hadoop. It stores large files by splitting them into blocks (default 128 MB) and distributing them across multiple nodes.
6	What is the Shuffle and Sort phase in MapReduce?	Shuffle and Sort is the intermediate phase between Map and Reduce. It groups all values with the same key together and sorts them before sending to the Reducer. It is handled automatically by Hadoop.
7	What is the difference between standalone and pseudo-distributed mode in Hadoop?	Standalone mode runs all Hadoop services in a single JVM using the local file system (no HDFS). Pseudo-distributed mode runs each service in a separate JVM on one machine and uses HDFS.
8	What is IntWritable and Text in Hadoop?	IntWritable and Text are Hadoop's serializable data types used instead of Java's int and String. They implement the Writable interface so Hadoop can serialize and transfer them across the network.

9	What is a Combiner in MapReduce?	A Combiner is a mini-reducer that runs on the mapper node before sending data to the reducer. It reduces network traffic by partially aggregating data locally. Example: summing word counts before shuffle.
10	What is the role of the Driver class in MapReduce?	The Driver class configures the MapReduce job. It sets the input/output paths, Mapper class, Reducer class, output key-value types and submits the job using <code>job.waitForCompletion()</code> .

Group B – Practical 2: MapReduce Log File Processing (Java)

Aim: *Design a distributed MapReduce application that processes a system log file to extract and aggregate information.*

Q.No	Question	Answer
1	What is log file processing?	Log file processing means analyzing system-generated log files to extract useful information like error counts, IP addresses, access times and response codes for monitoring and debugging.
2	How does MapReduce help in processing large log files?	MapReduce splits the large log file into chunks and processes them in parallel across multiple nodes. Each mapper processes one chunk and the reducer aggregates results, making it scalable.
3	What key-value pair does a Mapper emit for log processing?	For counting HTTP response codes, the Mapper emits (response_code, 1). Example: '404 Not Found' → (404, 1). The Reducer then sums counts for each response code.
4	What is the TextInputFormat in Hadoop?	TextInputFormat is the default input format in Hadoop. It reads input as lines of text and passes each line as a value (with byte offset as key) to the Mapper.
5	What is FileInputFormat and FileOutputFormat?	FileInputFormat.addInputPath(job, inputPath) sets the input directory. FileOutputFormat.setOutputPath(job, outputPath) sets the output directory where results are written on HDFS.
6	What is data skew in MapReduce?	Data skew means some reducers receive much more data than others, causing them to take longer. This slows down the entire job because MapReduce waits for all reducers to finish.
7	What are the steps to run a MapReduce job on Hadoop?	Steps: 1) Write Mapper, Reducer and Driver Java classes. 2) Compile and create a JAR file. 3) Copy input file to HDFS. 4) Run the JAR using 'hadoop jar'. 5) Check output on HDFS.
8	What command is used to put a file into HDFS?	hadoop fs -put localfile.txt /hdfs/path/ puts a local file onto HDFS. hadoop fs -get /hdfs/path output.txt retrieves a file from HDFS to local.
9	What is the use of context.write() in Hadoop?	context.write(key, value) is used inside Mapper and Reducer to emit output key-value pairs. In Mapper, it emits intermediate data. In Reducer, it writes the final output.

10	What command is used to check Hadoop services running?	The jps (Java Process Status) command lists all running Java processes including NameNode, DataNode, ResourceManager, NodeManager and JobHistoryServer.
----	--	---

Group B – Practical 3: Apache Spark with Scala

Aim: Write a simple program in Scala using the Apache Spark framework demonstrating RDD operations, transformations and actions.

Q.No	Question	Answer
1	What is Apache Spark?	Apache Spark is an open-source distributed computing framework for fast data processing. It processes data in memory (RAM), making it much faster than Hadoop MapReduce for iterative tasks.
2	What is the main difference between Hadoop MapReduce and Apache Spark?	Hadoop MapReduce stores intermediate data on disk after every step, making it slow. Apache Spark processes data in memory (RAM), making it 10–100x faster especially for iterative algorithms.
3	What is Scala?	Scala is a programming language that runs on the Java Virtual Machine (JVM). It combines object-oriented and functional programming. Apache Spark is written in Scala and has a native Scala API.
4	What is an RDD in Spark?	RDD (Resilient Distributed Dataset) is the fundamental data structure in Spark. It is a fault-tolerant, immutable, distributed collection of data that can be processed in parallel across a cluster.
5	What is the difference between transformation and action in Spark?	Transformations (map, filter, flatMap) are lazy – they define what to do but do not execute immediately. Actions (collect, count, reduce, saveAsTextFile) trigger actual execution.
6	What is lazy evaluation in Spark?	Lazy evaluation means Spark does not execute transformations immediately. It builds an execution plan (DAG) and only runs when an action is called. This allows Spark to optimize the plan.
7	What is SparkContext?	SparkContext is the entry point to any Spark application. It connects to the Spark cluster. In Scala: <code>val sc = new SparkContext(conf)</code> . All RDD operations start from SparkContext.
8	What is the difference between map() and flatMap() in Spark?	map() applies a function to each element and returns one output per input. flatMap() applies a function and returns zero or more outputs per input (flattens the result). Used in Word Count to split lines.

9	What is reduceByKey() in Spark?	reduceByKey(func) groups values by key and applies a reduce function. Example: <code>rdd.reduceByKey(_ + _)</code> sums all values for each key – used in Word Count to get total count per word.
10	How do you create an RDD from a text file in Spark (Scala)?	<code>val rdd = sc.textFile('input.txt')</code> . This creates an RDD where each element is one line of the file. You can then apply transformations like flatMap, map and reduceByKey.

